

Single Source Shortest Path

Dijkstra Algorithm

Single-Source Shortest Paths


Input: directed graph $G=(V, E)$. ($m=|E|$, $n=|V|$)

- each edge has non negative length l_e
- source vertex s

Output: for each $v \in V$, compute

$L(v) :=$ length of a shortest s - v path in G

Length of path
= sum of edge lengths




Assumption:

1. [for convenience] $\forall v \in V, \exists s \Rightarrow v$ path
2. [important] $l_e \geq 0 \quad \forall e \in E$


One of the following is the list of shortest---path distances for the nodes s, v, w, t , respectively. Which is it?

- a) 0,1,2,3
- b) 0,1,4,7
- c) 0,1,4,6
- d) 0,1,3,6




One of the following is the list of shortest---path distances for the nodes s, v, w, t , respectively. Which is it?

- a) 0,1,2,3
- b) 0,1,4,7
- c) 0,1,4,6
- d) 0,1,3,6



Why Another Shortest-Path Algorithm?

- Question: Doesn't BFS already compute shortest paths in linear time?
- Answer: yes, If $l_e = 1$ for every edge e .



- Question: why not just replace each edge e by directed path of l_e unit length edges:
- Answer: blows up graph too much
- Solution: Dijkstra's shortest path algorithm.



Dijkstra's Algorithm

This array
only to help
explanation!

Initialize:

- $X = [s]$ [vertices processed so far]
- $A[s] = 0$ [computed shortest path distances]
- $B[s] = \text{empty path}$ [computed shortest paths]

Main Loop

- while $X \neq V$:


-need to grow
x by one node

Main Loop cont'd:

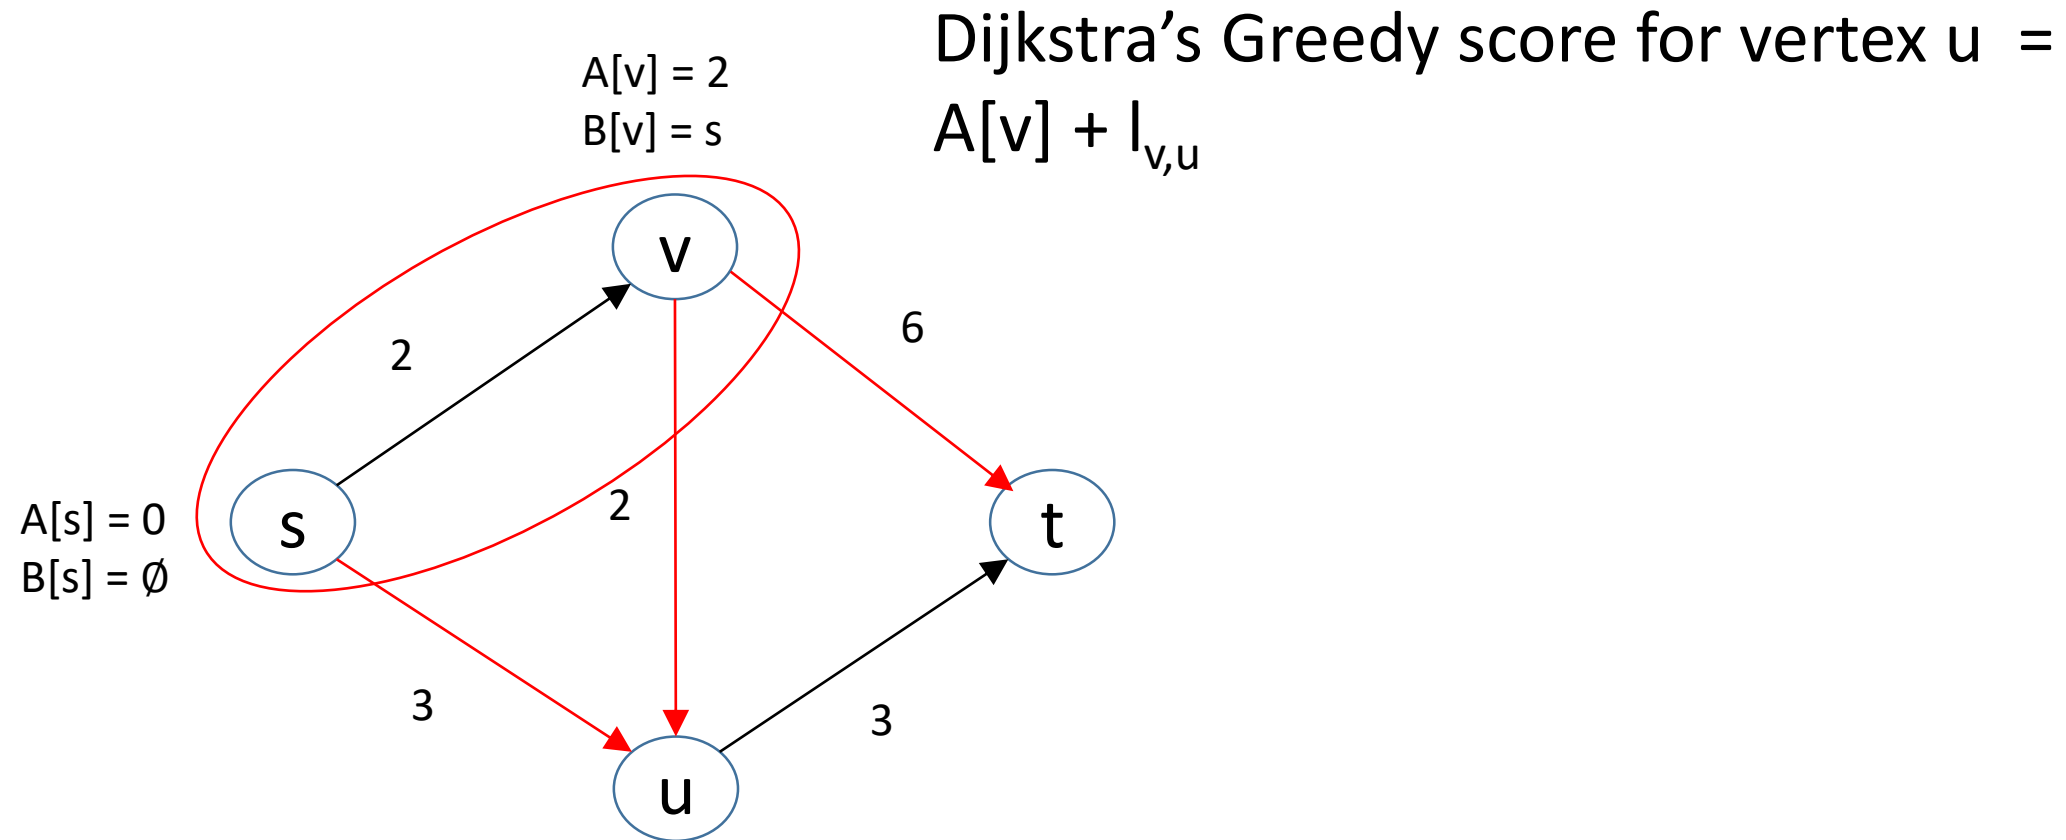
- among all edges $(v, w) \in E$
with $v \in X, w \notin X$,
pick the one that minimizes
 $A[v] + l_{vw}$
[call it (v^*, w^*)] → Already
computed in
earlier iteration
- add w^* to X
- set $A[w^*] := A[v^*] + l_{v^*w^*}$
- set $B[w^*] := B[v^*]u(v^*, w^*)$

Dijkstra's Algorithm

Dijkstra's Greedy score for vertex $u = A[v] + l_{v,u}$




Dijkstra's Algorithm

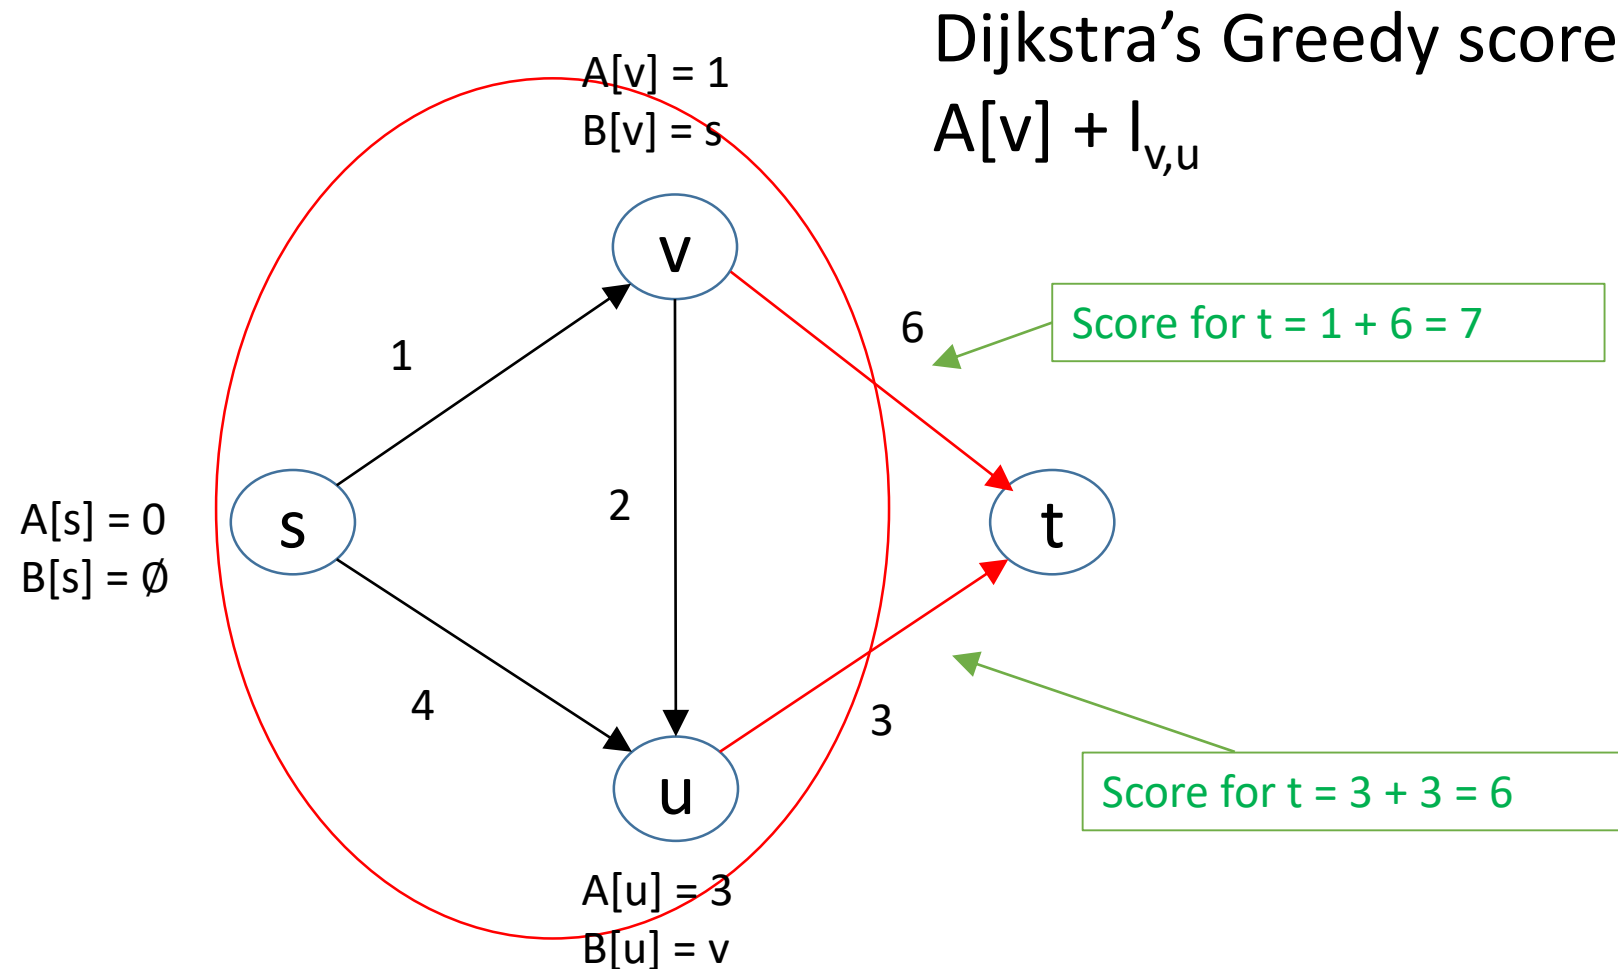


Dijkstra's Algorithm

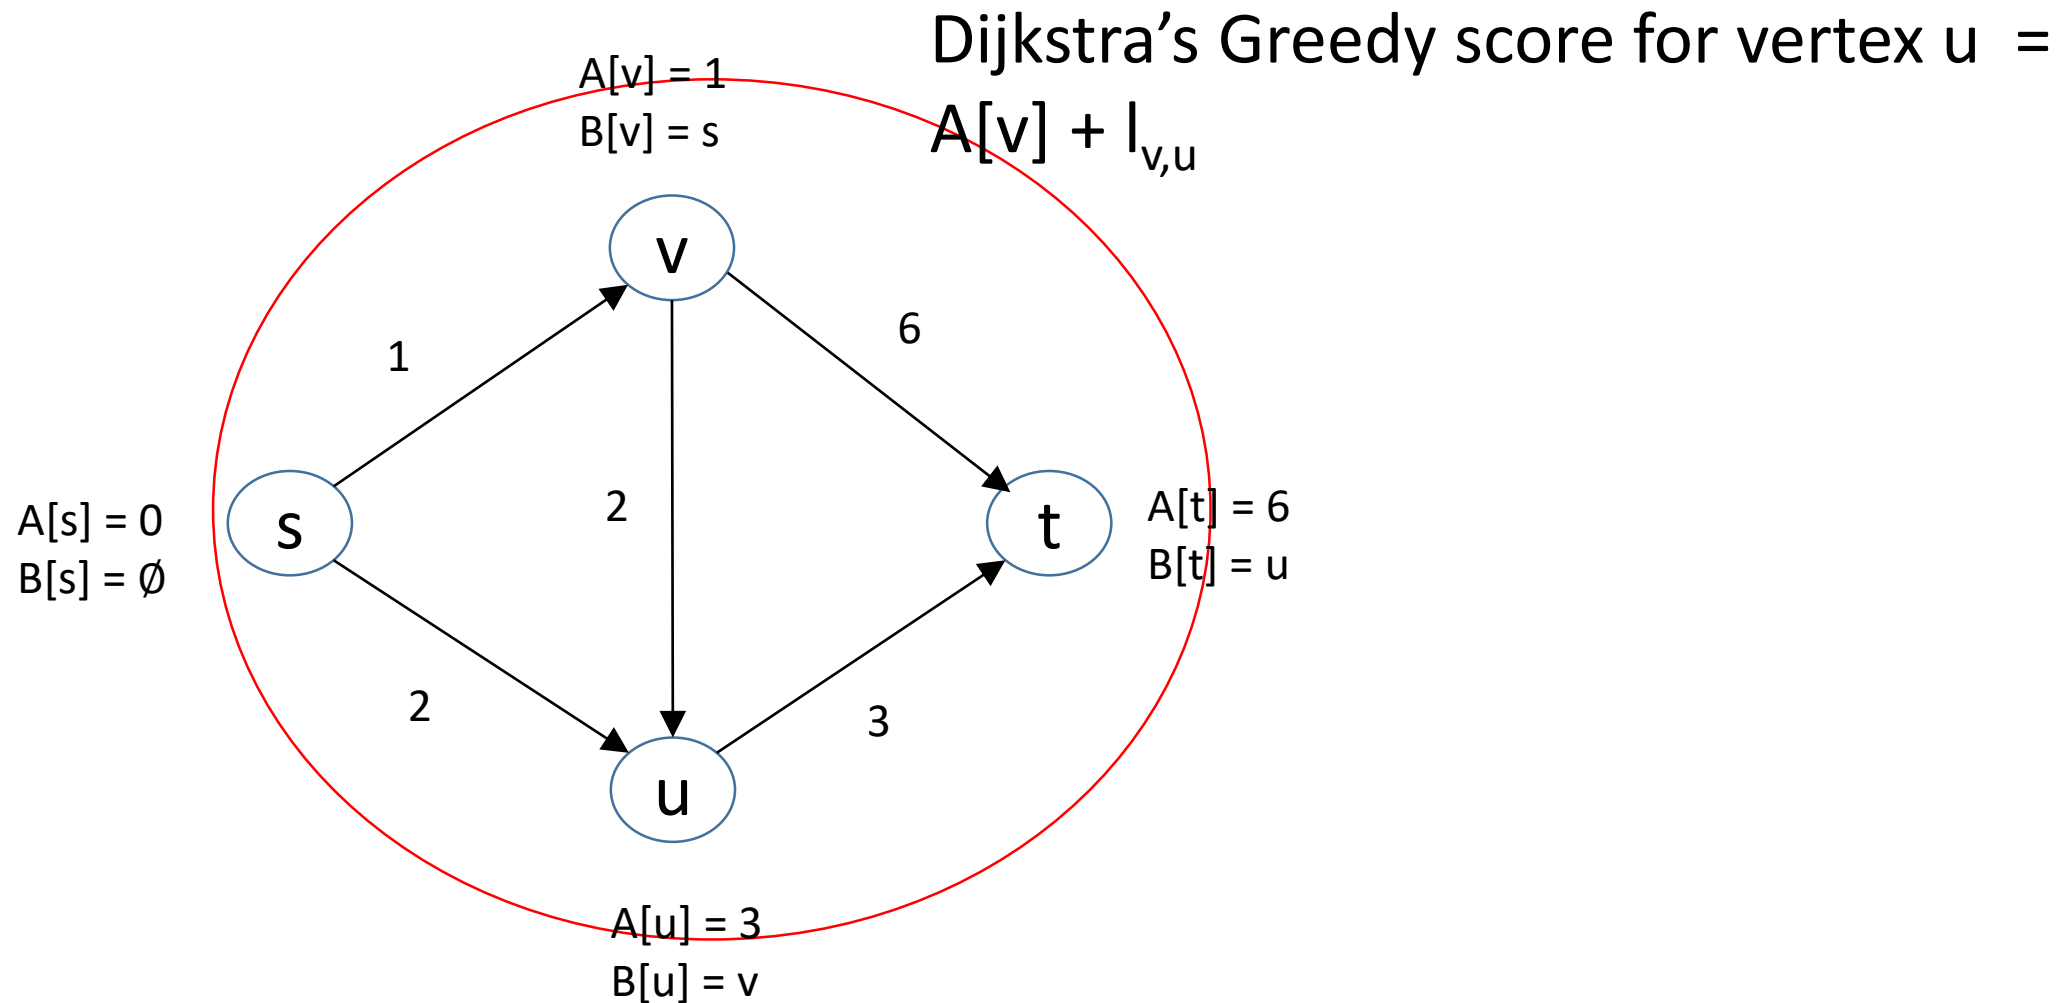
Dijkstra's Greedy score for vertex $u = A[v] + l_{v,u}$



Dijkstra's Algorithm




Dijkstra's Algorithm




Negative Edge Lengths

Dijkstra's algorithm incorrect on this graph !
(computes shortest s-t distance to be -2 rather than -4)




Negative Edge Lengths

Question: why not reduce computing shortest paths with negative edge lengths to the same problem with non negative lengths? (by adding large constant to edge lengths)




Negative Edge Lengths

Question: why not reduce computing shortest paths with negative edge lengths to the same problem with non negative lengths? (by adding large constant to edge lengths)



Negative Edge Lengths

- Question: why not reduce computing shortest paths with negative edge lengths to the same problem with non negative lengths? (by adding large constant to edge lengths)
- Problem: doesn't preserve shortest paths !
- Total weight added on path from s to t through direct edge is 5
- While total weight added on path from s to t through v is 10.



Dijkstra's Algorithm

This array
only to help
explanation!

Initialize:

- $X = [s]$ [vertices processed so far]
- $A[s] = 0$ [computed shortest path distances]
- $B[s] = \text{empty path}$ [computed shortest paths]

Main Loop

- while $X \neq V$:

-need to grow
x by one node

Main Loop cont'd:

- among all edges $(v, w) \in E$
with $v \in X, w \notin X$,
pick the one that minimizes
 $A[v] + l_{vw}$
[call it (v^*, w^*)] → Already
computed in
earlier iteration
- add w^* to X
- set $A[w^*] := A[v^*] + l_{v^*w^*}$
- set $B[w^*] := B[v^*]u(v^*, w^*)$

Dijkstra's Algorithm

- Which of the following running times seems to best describe a “naive” implementation of Dijkstra's algorithm?
 - a) $\theta(m+n)$
 - b) $\theta(m \log n)$
 - c) $\theta(n^2)$
 - d) $\theta(mn)$

Dijkstra's Algorithm

- Which of the following running times seems to best describe a “naive” implementation of Dijkstra's algorithm?

a) $\theta(m+n)$

b) $\theta(m \log n)$

c) $\theta(n^2)$

d) $\theta(mn)$

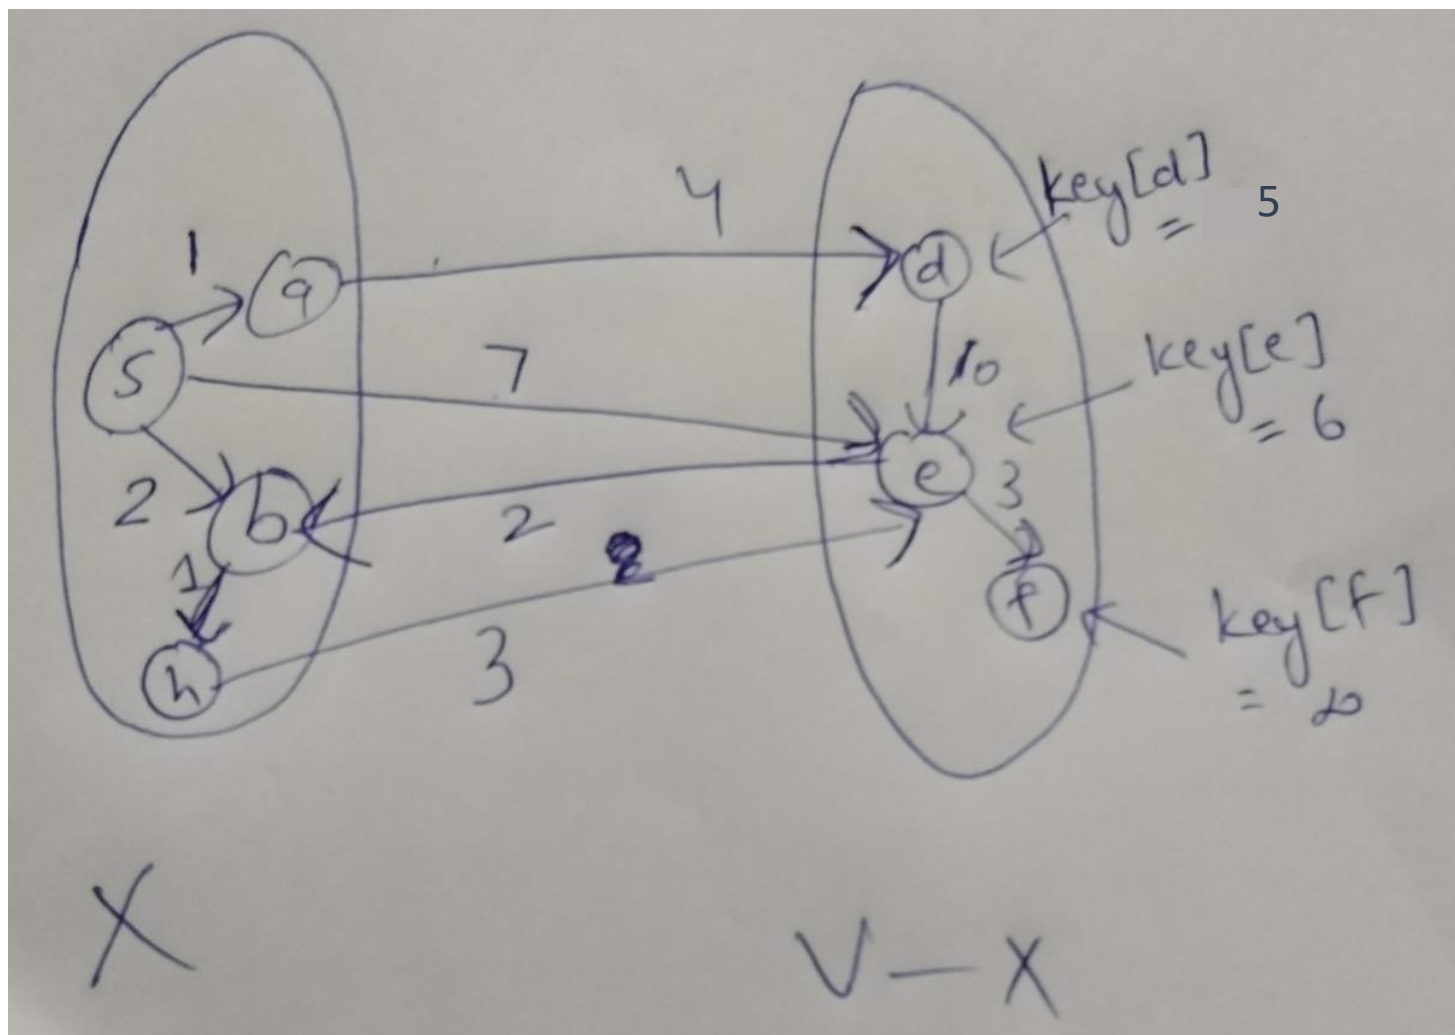
- $(n-1)$ iterations of while loop
- $O(m)$ work per iteration
- $O(1)$ [work per edge]

CAN WE DO BETTER?

Heap Operations

- perform Insert, Extract-Min, delete in $O(\log n)$ time.
- conceptually, a perfectly balanced binary tree
- Heap property: at every node, $\text{key} \leq \text{children's keys}$
- extract-min by swapping up last leaf, bubbling down
- insert via bubbling up
- Also: will need ability to delete from middle of heap. (bubble up or down as needed)





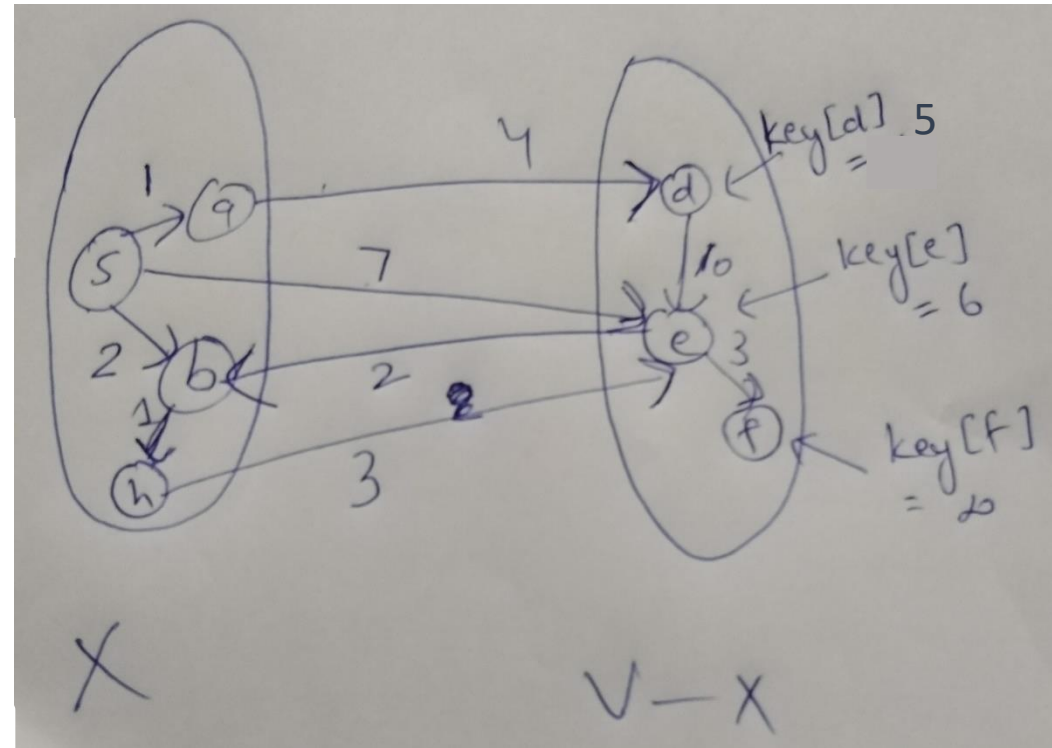
Two Invariants

Invariant # 1: elements in heap = vertices of $V-X$.

Invariant #2: for $v \notin X$

$\text{Key}[v]$ = smallest Dijkstra greedy score of an edge (u, v) in E with u in X

(of $+\infty$ if no such edges exist)



Dijkstra's greedy score of (v, w) : $A[v] + l_{vw}$

Point: by invariants, Extract-Min yields correct vertex w^* to add to X next.

(and we set $A[w^*]$ to $\text{key}[w^*]$)

Maintaining the Invariants

To maintain Invariant #2: [i.e., that $\forall v \notin X$

Key[v] = smallest Dijkstra greedy score of edge (u,v) with u in X]

When w extracted from heap (i.e., added to X)


- for each edge (w,v) in E:
 - if v in V-X (i.e., in heap)

Key update {


- delete v from heap
- recompute $\text{key}[v] = \min \{ \text{key}[v], A[w] + l_{wv} \}$
- re-Insert v into heap

Greedy score of (w,v)


Dijkstra Algorithm



Dijkstra Algorithm



Dijkstra Algorithm



Dijkstra's Algorithm

```
Dijkstra(G,w,s)
{
    % Initialize
    for (each  $u \in V$ )
    {
         $d[u] = \infty$ ;
         $color[u] = \text{white}$ ;
    }
     $d[s] = 0$ ;
     $pred[s] = \text{NIL}$ ;
     $Q = (\text{queue with all vertices})$ ;

    while (Non-Empty( $Q$ ))
    {
        % Process all vertices
         $u = \text{Extract-Min}(Q)$ ;
        % Find new vertex
        for (each  $v \in Adj[u]$ )
        {
            % If estimate improves
            if ( $d[u] + w(u, v) < d[v]$ )
            {
                 $d[v] = d[u] + w(u, v)$ ;
                % relax
                Decrease-Key( $Q, v, d[v]$ );
                 $pred[v] = u$ ;
            }
        }
         $color[u] = \text{black}$ ;
    }
}
```

Running Time Analysis


You check: dominated by heap operations. ($O(\log(n))$ each)


- $(n-1)$ Extract mins
- each edge (v,w) triggers at most one Delete/Insert combo (if v added to X first)

So: # of heap operations in $O(n+m) = O(m)$

So: running time = $O(m \log(n))$ (like sorting)

Since graph is
weakly connected






Step 0: Initialization.

v	s	a	b	c	d
$d[v]$	0	∞	∞	∞	∞
$pred[v]$	nil	nil	nil	nil	nil
$color[v]$	W	W	W	W	W

Priority Queue:

v	s	a	b	c	d
$d[v]$	0	∞	∞	∞	∞




Step 1: As $Adj[s] = \{a, b\}$, work on a and b and update information.

v	s	a	b	c	d
$d[v]$	0	2	7	∞	∞
$pred[v]$	nil	s	s	nil	nil
$color[v]$	B	W	W	W	W

Priority Queue:

v	a	b	c	d
$d[v]$	2	7	∞	∞




Step 2: After Step 1, a has the minimum key in the priority queue. As $Adj[a] = \{b, c, d\}$, work on b, c, d and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	10	7
$pred[v]$	nil	s	a	a	a
$color[v]$	B	B	W	W	W

Priority Queue:

v	b	c	d
$d[v]$	5	10	7




Step 3: After Step 2, b has the minimum key in the priority queue. As $Adj[b] = \{a, c\}$, work on a, c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	W	W

Priority Queue:

v	c	d
$d[v]$	6	7




Step 4: After Step 3, c has the minimum key in the priority queue. As $Adj[c] = \{d\}$, work on d and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	W

Priority Queue:

v	d
$d[v]$	7



Step 5: After Step 4, d has the minimum key in the priority queue. As $Adj[d] = \{c\}$, work on c and update information.

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a
$color[v]$	B	B	B	B	B


Priority Queue: $Q = \emptyset$.

We are done.

Shortest Path Tree: $T = (V, A)$, where

$$A = \{(pred[v], v) | v \in V \setminus \{s\}\}.$$

The array $pred[v]$ is used to build the tree.



Example:

v	s	a	b	c	d
$d[v]$	0	2	5	6	7
$pred[v]$	nil	s	a	b	a

Correctness Claim

Theorem [Dijkstra] For every directed graph with nonnegative edge lengths, Dijkstra's algorithm correctly computes all shortest-path distances.

$$[i.e., A[v] = L(v) \quad \forall v \in V]$$

what algorithm
computes

True shortest
distance from s to v

Proof: by induction on the number of iterations.

Base Case: $A[s] = L[s] = 0$ (correct)

Proof

Inductive Step:

Inductive Hypothesis: all previous iterations correct (i.e., $A[v] = L(v)$ and $B[v]$ is a true shortest s - v path in G , for all v already in X).

In current iteration:

We pick an edge (v^*, w^*) and we add w^* to X .


We set $B[w^*] = B[v^*] \cup (v^*, w^*)$

has length $L(v^*) + l_{v^*w^*}$

has length $L(v^*)$

$L(v^*)$ by I.H

Also: $A[w^*] = A[v^*] + l_{v^*w^*} = L(v^*) + l_{v^*w^*}$



Proof (con'd)


Upshot: in current iteration, we set:

1. $A[w^*] = L(v^*) + l_{v^*w^*}$
2. $B[w^*] = \text{an } s \rightarrow w^* \text{ path with length } (L(v^*) + l_{v^*w^*})$

To finish proof: need to show that *every* $s \rightarrow w^*$ path has length \geq
 $L(v^*) + l_{v^*w^*}$ (if so, our path is the shortest!)


So: Let $P = \text{any } s \rightarrow w^*$ path. Must “cross the frontier”:

and so has the form:



Proof (con'd)

So: every $s \rightarrow w^*$ path P has to have the form



Total length of path P: at least $A[y] + C_{yz}$ length of our path !

\rightarrow by Dijkstra's greedy criterion, $A[v^*] + l_{v^*w^*} \leq A[y] + l_{yz} \leq \text{length of } P$

($y \in X$
 $z \notin X$)