


Introduction to **Machine Learning** & Linear Regression – **Model Representation**

Acknowledgments: Andrew Ng and Pedro Domingos


So What Is Machine Learning?

- ❑ Automating automation
- ❑ Getting computers to program themselves
- ❑ Let the **data** do the work!

□ Traditional Programming



□ Machine Learning



Machine Learning

- Grew out of work in AI
- New capability for computers

Examples:


- Database mining
 - Large datasets from growth of automation/web.
 - E.g., Web click data, medical records, biology, engineering
- Applications can't program by hand.
 - E.g., Autonomous helicopter, handwriting recognition, most of Natural Language Processing (NLP), Computer Vision.

Practical Applications of Machine Learning

- Spam filtering
- Speech/handwriting recognition
- Object detection/recognition
- Weather prediction
- Stock market analysis
- Search engines (e.g, Google)
- Ad placement on websites
- Credit-card fraud detection
- Webpage clustering (e.g., Google News)
- Social Network Analysis
- Recommendation systems (e.g., Netflix, Amazon)
- Automatic vehicle navigation

Machine Learning definition

- Arthur Samuel (1959). Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.



- Tom Mitchell (1998) Well-posed Learning Problem: A computer program is said to *learn* from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience E .

“A computer program is said to *learn* from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience E .”

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task T in this setting?

- a) Watching you label emails as spam or not spam.
- b) Classifying emails as spam or not spam.
- c) The number (or fraction) of emails correctly classified as spam/not spam.
- d) None of the above—this is not a machine learning problem.


Machine learning algorithms:

- Supervised learning
- Unsupervised learning

Others: recommender systems.

Introduction to **Supervised Learning**

Housing price prediction.




Supervised Learning

“right answers” given

Regression: Predict continuous


valued output (price)

Cancer (malignant, benign)



Classification

Discrete valued
output (0 or 1)



- Clump Thickness
- Uniformity of Cell Size
- Uniformity of Cell Shape
- ...

You're running a company, and you want to develop learning algorithms to address each of two problems.

Problem 1: You have a large inventory of identical items. You want to predict how many of these items will sell over the next 3 months.


Problem 2: You'd like software to examine individual customer accounts, and for each account decide if it has been hacked/compromised.

Should you treat these as classification or as regression problems?


- a) Treat both as classification problems.
- b) Treat problem 1 as a classification problem, problem 2 as a regression problem.
- c) Treat problem 1 as a regression problem, problem 2 as a classification problem.
- d) Treat both as regression problems.

Introduction to **Unsupervised Learning**

Supervised Learning



Unsupervised Learning



Google News

news.google.com

Web Images Videos Maps News Shopping Gmail more

Search News

Search the Web

Advanced news search

U.S. edition Add a section

Top Stories

Deepwater Horizon
Fed meeting
Foreign exchange market
Lindsay Lohan
iGht
Tom Brady
Toronto
International Film Festival
Paris Hilton
Iran
Hurricane Igor

Top Stories

Christine O'Donnell
White House official denies Tea Party-focused ad campaign
CNN International - Ed Henry - 1 hour ago
Democratic sources say the White House is not considering an ad campaign tying Republicans to the Tea Party. Washington (CNN) - A top White House official sharply denied a report that claims President Obama's political advisers are weighing a national...
Tea Party is replacing the blame, former President Bill Clinton claims
New York Daily News
GOP tea party backer defends Christine O'Donnell The Associated Press
Atlanta Journal Constitution - Politics Daily - MyFox Washington DC - Salon
all 276 news articles

US Stocks Climb After Recession Called Over, Homebuilders Gain
MarketWatch - Kristina Peterson - 16 minutes ago
NEW YORK (MarketWatch) - US stocks climbed Monday, gaining speed after a key nonprofit organization officially called the recession over, giving investors a boost of confidence in the gradual economic recovery.
Longest recession since 1930s ended in June 2009, group says
Los Angeles Times
Downturn Was Longest in Decades, Panel Confirms New York Times
Wall Street Journal - AFP - CNN - USA Today
all 276 news articles

BP Oil Well, Site of National Catastrophe, Dies at One
Vanity Fair - Juli Weiner - 22 minutes ago
The BP oil well, site of the Deepwater Horizon explosion that led to the worst oil spill in US history, died today at one year old.
Video: Blown-out BP Well Finally Killed in Gulf of Mexico The Associated Press
Weiss Doubts BP Would End Operations in Gulf of Mexico Video Blogging
CNN International - Wall Street Journal (blog) - The Guardian -
all 2,292 news articles

Recession
CNN Money - Chris Isidore - 39 minutes ago
Hurricane Igor lashes Bermuda USA Today - Gerry Boone - 5 minutes ago
Explain what you want from us reads front-page editorial msnbc.com - Olivia Torres - 10 minutes ago

Crisis response: Pakistan floods
San Francisco Bay Area - Edit
Clerx
Bay Biz Buzz: Clerx close to selling STP, Amcor All
San Jose Mercury News - 48 minutes ago
all 24 articles
Google's official browser keeps the company buzzing msnbc.com
San Jose Mercury News - Bruce Newman - 1 hour ago
Jan Sylvia
Martinez man still unconscious as police investigate weekend shooting
San Jose Mercury News - Robert Salanga - 48 minutes ago - all 6 articles

Spotlight
Sarkozy rages at EU 'humiliation'
Financial Times - Peggy Hollinger - September 20, 2010

BP Kills Macondo, But Its...
blogs.wsj.com/source/2010/09/20/bp-kills-macondo-but-its-legacy-lives-on/

THE WALL STREET JOURNAL

Log In • Register for Free • Subscribe Now, Get 2 Weeks Free

THE SOURCE

Financial Services Transport Leisure Insurance Oil & Gas Sport Caught on the Web Betting Technology

SEPTEMBER 20, 2010, 12:44 PM GMT

BP Kills Macondo, But Its Legacy Lives On

Search The Source

Article Comments (2)

THE SOURCE HOME PAGE

By James Herron

BP confirmed late Sunday that the Macondo well that leaked almost five million barrels of oil into the Gulf of Mexico has been permanently sealed, but the well will continue to affect BP and the wider oil industry for many years.

The most immediate worry for BP and its shareholders is how the authorities will apportion blame for the spill. BP's own investigation corded responsibility across

About The Source

The Source is WSJ.com Europe's home for rapid-fire analysis of the day's big business and finance stories. It is edited by Lauren Mills, based in London.

Most Recent

Articles Comments

- Who Needs Plaza II Anyway
- Will Banks Be Forced to Split Retail And Banking Arms?
- Timing of Ratings Award Intriguing
- BP Kills Macondo, But Its Legacy Lives On
- We Already Need a Second to Brawl III

Allen: Well is dead, but m...
edition.cnn.com/10/US/09/20/gulf.oil.disaster/

EDITION: INTERNATIONAL U.S. MEXICO ARABIC

Home Video World U.S. Africa Asia Europe Latin America Middle East Business

Allen: Well is dead, but much Gulf Coast work remains

By the CNN Wire Staff

September 20, 2010 - Updated 1317 GMT (2117 HKT)

Click to play

What next for Gulf oil spill?

STORY HIGHLIGHTS

(CNN) -- The ruptured Macondo well, a mile under the Gulf of Mexico off the Louisiana coast, has been nonpermanently dead.

BP oil spill cost hits nearl...
www.guardian.co.uk/environment/2010/sep/20/bp-oil-spill-dee

Mobile site Sign in Register Text larger smaller About us

guardian.co.uk

News Sport Comment Culture Business Money Life & style


Business BP

BP oil spill cost hits nearly \$10bn



BP has set up a \$20bn compensation fund after the Deepwater Horizon disaster, which has so far paid out 19,000 claims totalling more than \$240m

Julia Kollewe
guardian.co.uk Monday 20 September 2010 08:33 BST
Article history

BP's costs for the Deepwater Horizon disaster have hit \$10bn. Photograph: HoReuters



Organize computing clusters



Social network analysis



Astronomical data analysis

Of the following examples, which would you address using an unsupervised learning algorithm? (Check all that apply.)

Given email labeled as spam/not spam, learn a spam filter.

Given a set of news articles found on the web, group them into set of articles about the same story.

Given a database of customer data, automatically discover market segments and group customers into different market segments.


Given a dataset of patients diagnosed as either having diabetes or not, learn to classify new patients as having diabetes or not.

Linear regression
with one variable

Model representation

Housing Prices

Price
(in 1000s
of dollars)



Supervised Learning

Given the “right answer” for each example in the data.

Regression Problem

Predict real-valued output

Training set of housing prices	Size in feet ² (x)	Price (\$) in 1000's (y)
	2104	460
	1416	232
	1534	315
	852	178

Notation:

m = Number of training examples

x's = “input” variable / features

y's = “output” variable / “target” variable

Training Set



Learning Algorithm



Size of
house



h



Estimated
price

How do we represent h ?

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Linear regression with one variable.
Univariate linear regression.

Linear regression
with one variable

Cost function


Training Set	Size in feet ² (x)	Price (\$) in 1000's (y)
	2104	460
	1416	232
	1534	315
	852	178

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

θ_i 's: Parameters


How to choose θ_i 's ?

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$




$$\theta_0 = 1.5$$

$$\theta_1 = 0$$




$$\theta_0 = 0$$

$$\theta_1 = 0.5$$



$$\theta_0 = 1$$

$$\theta_1 = 0.5$$



Cost Function
(Squared Error function):

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

Idea: Choose θ_0, θ_1 so that
 $h_{\theta}(x)$ is close to y for our
training examples (x, y)

Linear regression
with one variable

Cost function
intuition I

Hypothesis:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Parameters:

$$\theta_0, \theta_1$$

Cost Function:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

Simplified (to understand
cost function better)

$$h_{\theta}(x) = \theta_1 x$$


$$\theta_1$$

$$J(\theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

minimize $J(\theta_1)$
 θ_1


$$h_{\theta}(x)$$

(for fixed $\theta_1=1$, this is a function of x)




$$J(\theta_1)$$

(function of the parameter $\theta_1=1$)




$$h_{\theta}(x)$$

(for fixed $\theta_1=1$, this is a function of x)




$$J(\theta_1)$$

(function of the parameter $\theta_1=1$)




$$h_{\theta}(x)$$

(for fixed $\theta_1 = 0.5$, this is a function of x)




$$J(\theta_1)$$

(function of the parameter θ_1)




$$h_{\theta}(x)$$

(for fixed $\theta_1 = 0$, this is a function of x)



$$J(\theta_1)$$

(function of the parameter θ_1)



Linear regression
with one variable

Cost function
intuition II

Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$


Parameters: θ_0, θ_1

Cost Function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

$$h_{\theta}(x)$$


(for fixed θ_0, θ_1 , this is a function of x)



$$\theta_0 = 50 \quad \theta_1 = 0.06$$

$$J(\theta_0, \theta_1)$$


(function of the parameters θ_0, θ_1)



$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

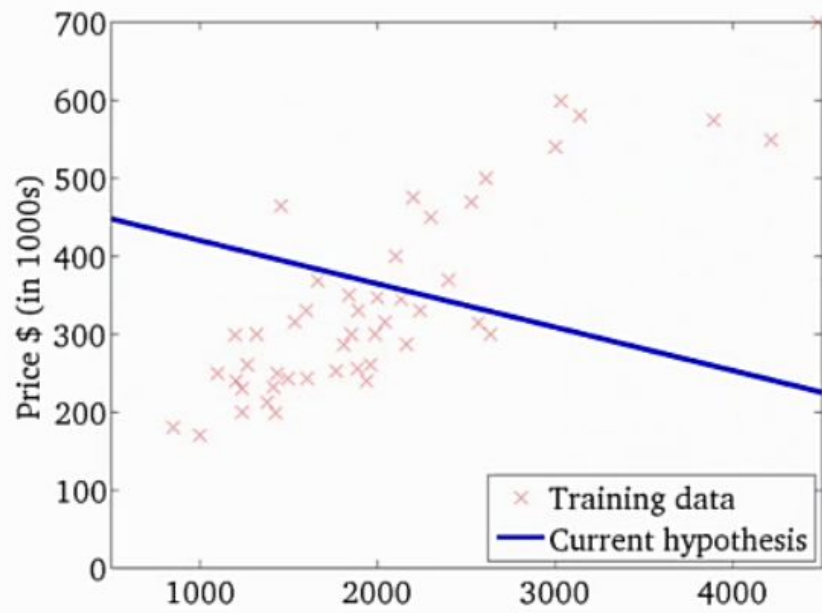
$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



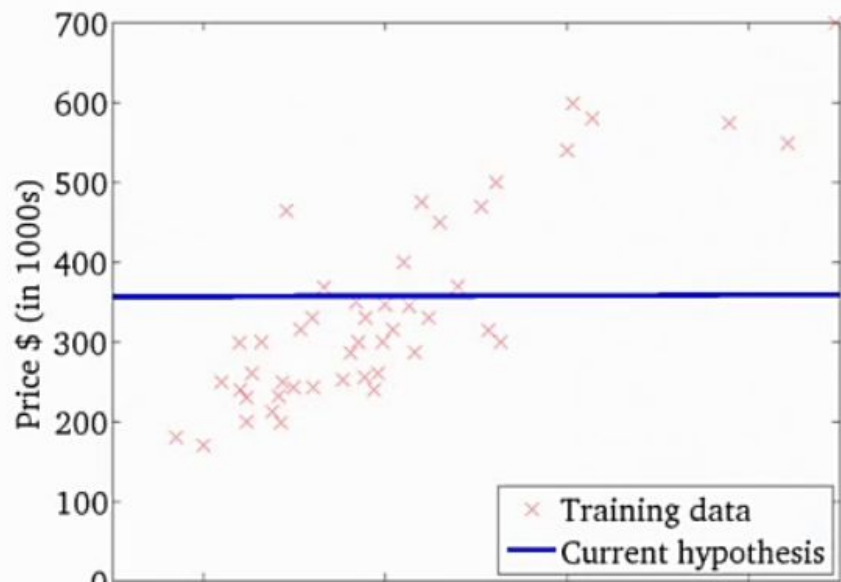
$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)




$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)




$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)



$$h_{\theta}(x)$$

(for fixed θ_0, θ_1 , this is a function of x)



$$J(\theta_0, \theta_1)$$

(function of the parameters θ_0, θ_1)

