

Parallel and Distributed Computing

CS3006 (BDS-6A)

Lecture 07


Instructor: Dr. Syed Mohammad Irteza

Assistant Professor, Department of Computer Science, FAST


16 February, 2023

Principles of Parallel Algorithm Design

Task Interact Graph (Sparse-matrix multiplication)



(a)



(b)

Figure 3.6 A decomposition for sparse matrix-vector multiplication and the corresponding task-interaction graph. In the decomposition Task i computes $\sum_{0 \leq j \leq 11, A[i,j] \neq 0} A[i,j] \cdot b[j]$.

Processes and Mapping

- The *logical processing* or *computing agent* that performs tasks is called a **process**.
- The *mechanism* by which tasks are *assigned to processes* for execution is called **mapping**.
- *Multiple tasks* can be mapped on a *single process*
- *Independent tasks* should be mapped onto *different processes*
- *Tasks* with high *mutual-interactions* should be mapped onto a *single process*

Processes and Mapping





Figure 3.7 Mappings of the task graphs of Figure 3.5 onto four processes.

Processes vs. Processors

- **Processes** are logical computing agents that perform tasks
- **Processors** are the hardware units that physically perform computations
- Depending on the problem, *multiple processes* can be mapped onto a *single processor*
- But, in most of the cases, there is a *one-to-one correspondence between processors and processes*
- So, we assume that there are *as many processes* as the number of *physical CPUs* on the parallel computer

Exercise:

- For the task graph given above, determine:
- Maximum degree of concurrency
- Critical path
- Maximum possible speedup assuming large number of process are available
- Minimum number of processes needed to obtain the maximum possible speedup
- Maximum speedup if number of processes are limited to 4



ASSUME: All weights are one, and arcs are directing upwards. We may assume that average degree of concurrency can be used as the maximum possible speedup.

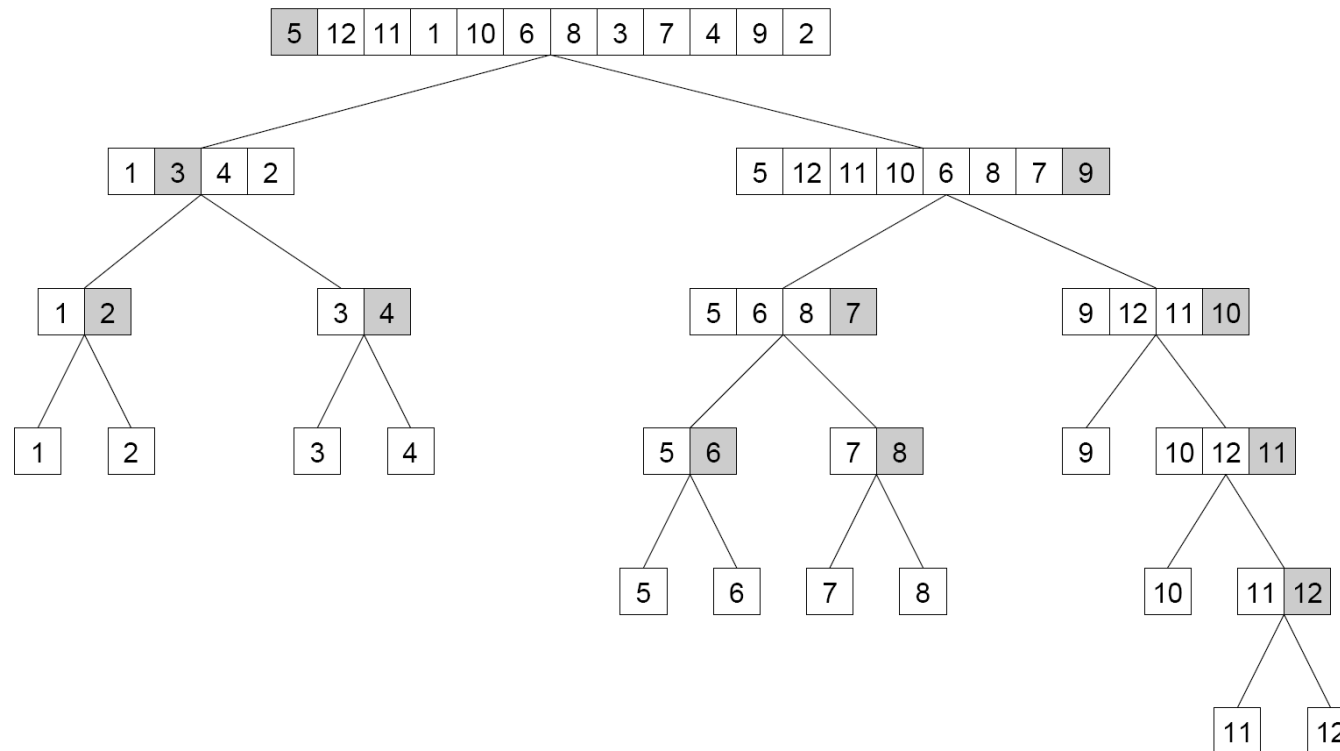
Decomposition Techniques

- The process of decomposing larger problems into smaller tasks for concurrent executions is known as *decomposition*.
- The techniques that facilitate this decomposition are known as *decomposition techniques*.
- Common techniques:
 - Recursive
 - Data-decomposition
 - Exploratory decomposition
 - Speculative decomposition
 - Hybrid
- *Recursive* and *data decompositions* are relatively *general purpose*
- *Exploratory* and *speculative* are *special purpose* in nature

Recursive Task Decomposition

- Recursive decomposition is a method for inducing concurrency in the problems that can be solved using a *divide and conquer strategy*
- It *divides* each problem into a *set of independent sub-problems*
- Each one of these *subproblems* is solved by *recursively applying* a similar *division* into *smaller subproblems* followed by a *combination of their results*
- A *natural concurrency* exists as *different subproblems* can be *solved concurrently*.

Recursive Decomposition



At each level and for each vector

1. Select a pivot
2. Partition set around pivot
3. Recursively sort each subvector




Figure 3.8 The quicksort task-dependency graph based on recursive decomposition for sorting a sequence of 12 numbers.

Recursive Decomposition

Modifying simple problem to
support recursive decomposition

```
1.  procedure SERIAL_MIN ( $A, n$ )
2.  begin
3.     $min = A[0]$ ;
4.    for  $i := 1$  to  $n - 1$  do
5.      if ( $A[i] < min$ )  $min := A[i]$ ;
6.    endfor;
7.    return  $min$ ;
8.  end SERIAL_MIN
```

Algorithm 3.1 A serial program for finding the minimum in an array of numbers A of length n .

```
1.  procedure RECURSIVE_MIN ( $A, n$ )
2.  begin
3.    if ( $n = 1$ ) then
4.       $min := A[0]$ ;
5.    else
6.       $lmin := RECURSIVE\_MIN (A, n/2)$ ;
7.       $rmin := RECURSIVE\_MIN (\&(A[n/2]), n - n/2)$ ;
8.      if ( $lmin < rmin$ ) then
9.         $min := lmin$ ;
10.     else
11.        $min := rmin$ ;
12.     endelse;
13.   endelse;
14.   return  $min$ ;
15. end RECURSIVE_MIN
```

Algorithm 3.2 A recursive program for finding the minimum in an array of numbers A of length n .

Recursive Decomposition

Modifying simple problem to support recursive decomposition




Figure 3.9 The task-dependency graph for finding the minimum number in the sequence {4, 9, 1, 7, 8, 11, 2, 12}. Each node in the tree represents the task of finding the minimum of a pair of numbers.

Data Decomposition

- Powerful and commonly used method
- Two step procedure:
 - *Partition data* on which computation is to be performed
 - This data partitioning is used to *induce* a *partitioning of the computations* into tasks.
- Partitioning output data
 - Used where each element of the output *can be computed independently* of others *as a function of the input*.
 - Partitioning of the output data *automatically induces a decomposition* of the problems into tasks
 - Each task is assigned *the work of computing a portion of the output*

Data Decomposition (Partitioning Output Data)

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

(a)

Task 1: $C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1}$

Task 2: $C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2}$

Task 3: $C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1}$

Task 4: $C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2}$

(b)

Figure 3.10 (a) Partitioning of input and output matrices into 2×2 submatrices. (b) A decomposition of matrix multiplication into four tasks based on the partitioning of the matrices in (a).

Data Decomposition (Partitioning Output Data)

Decomposition I	Decomposition II
Task 1: $C_{1,1} = A_{1,1}B_{1,1}$	Task 1: $C_{1,1} = A_{1,1}B_{1,1}$
Task 2: $C_{1,1} = C_{1,1} + A_{1,2}B_{2,1}$	Task 2: $C_{1,1} = C_{1,1} + A_{1,2}B_{2,1}$
Task 3: $C_{1,2} = A_{1,1}B_{1,2}$	Task 3: $C_{1,2} = A_{1,2}B_{2,2}$
Task 4: $C_{1,2} = C_{1,2} + A_{1,2}B_{2,2}$	Task 4: $C_{1,2} = C_{1,2} + A_{1,1}B_{1,2}$
Task 5: $C_{2,1} = A_{2,1}B_{1,1}$	Task 5: $C_{2,1} = A_{2,2}B_{2,1}$
Task 6: $C_{2,1} = C_{2,1} + A_{2,2}B_{2,1}$	Task 6: $C_{2,1} = C_{2,1} + A_{2,1}B_{1,1}$
Task 7: $C_{2,2} = A_{2,1}B_{1,2}$	Task 7: $C_{2,2} = A_{2,1}B_{1,2}$
Task 8: $C_{2,2} = C_{2,2} + A_{2,2}B_{2,2}$	Task 8: $C_{2,2} = C_{2,2} + A_{2,2}B_{2,2}$

Figure 3.11 Two examples of decomposition of matrix multiplication into eight tasks.

Data Decomposition (Partitioning Output Data)

(a) Transactions (input), itemsets (input), and frequencies (output)

Database Transactions	A, B, C, E, G, H	Itemsets	A, B, C	Itemset Frequency	1
	B, D, E, F, K, L		D, E		3
	A, B, F, H, L		C, F, G		0
	D, E, F, H		A, E		2
	F, G, H, K,		C, D		1
	A, E, F, K, L		D, K		2
	B, C, D, G, H, L		B, C, F		0
	G, H, L		C, D, K		0
	D, E, F, K, L				
	F, G, H, L				

(b) Partitioning the frequencies (and itemsets) among the tasks

Database Transactions	A, B, C, E, G, H	Itemsets	A, B, C	Itemset Frequency	1
	B, D, E, F, K, L		D, E		3
	A, B, F, H, L		C, F, G		0
	D, E, F, H		A, E		2
	F, G, H, K,				
	A, E, F, K, L				
	B, C, D, G, H, L				
	G, H, L				
	D, E, F, K, L				
	F, G, H, L				

task 1

Database Transactions	A, B, C, E, G, H	Itemsets	C, D	Itemset Frequency	1
	B, D, E, F, K, L		D, K		2
	A, B, F, H, L		B, C, F		0
	D, E, F, H		C, D, K		0
	F, G, H, K,				
	A, E, F, K, L				
	B, C, D, G, H, L				
	G, H, L				
	D, E, F, K, L				
	F, G, H, L				

task 2

Figure 3.12 Computing itemset frequencies in a transaction database.

Decomposition Techniques

Data Decomposition

Partitioning *input* data

- In many algorithms, it is not possible or desirable to *partition the output data*.
 - The output may be a *single unknown value*.
 - Such as in case of *finding sum, minimum, maximum* or *frequencies of a number*.
- It is sometimes possible to *partition the input data*, and then use this partitioning to *induce concurrency*
- A *task is created for each partition of the input data* and this task performs *as much computation as possible using this local data*
- Then *local solutions are combined* to generate a *global solution*

Decomposition Techniques

Partitioning input data

(a) Partitioning the transactions among the tasks

Database Transactions	A, B, C, E, G, H	Itemsets	A, B, C	Itemset Frequency	1
	B, D, E, F, K, L		D, E		2
	A, B, F, H, L		C, F, G		0
	D, E, F, H		A, E		1
	F, G, H, K,		C, D		0
			D, K		1
			B, C, F		0
			C, D, K		0
task 1					

Database Transactions		Itemsets	A, B, C	Itemset Frequency	0
			D, E		1
			C, F, G		0
	A, E, F, K, L		A, E		1
	B, C, D, G, H, L		C, D		1
	G, H, L		D, K		1
	D, E, F, K, L		B, C, F		0
	F, G, H, L		C, D, K		0
task 2					

Decomposition Techniques

Data Decomposition

Partitioning *both input and output* data

- Consider the problems where output data-partitioning is possible
- Here, partitioning the input also, can offer additional concurrency
- The next example shows *4-way decomposition of the previous example* based on both input-output partitioning.

Decomposition Techniques

Partitioning *both input and output data*

(b) Partitioning both transactions and frequencies among the tasks

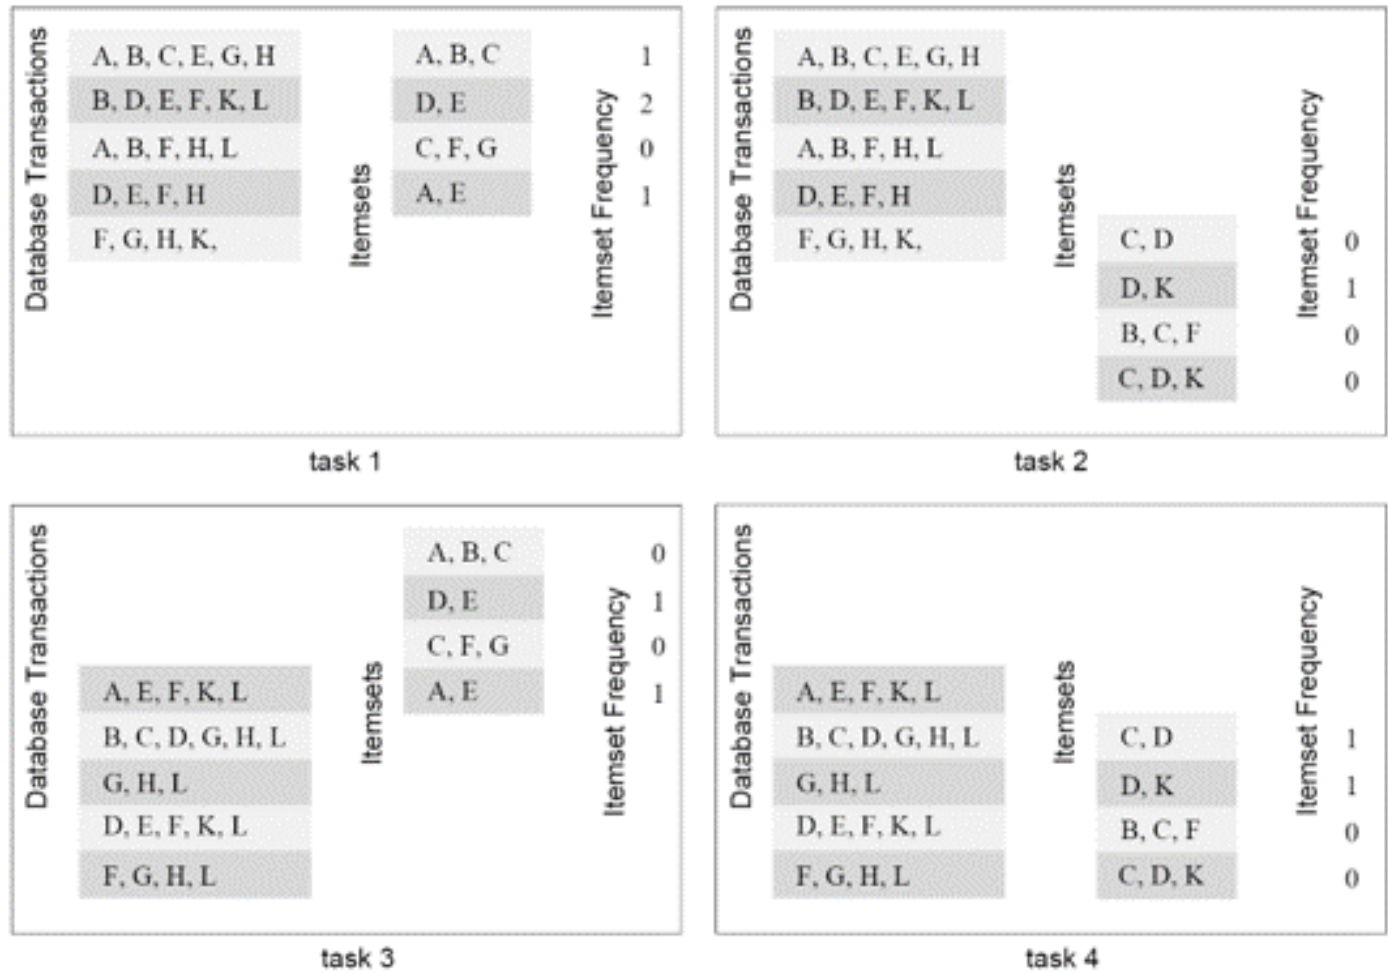


Figure 3.13 Some decompositions for computing itemset frequencies in a transaction database.

Decomposition Techniques

Partitioning *both intermediate data*

Stage I

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \left(\begin{pmatrix} D_{1,1,1} & D_{1,1,2} \\ D_{1,2,2} & D_{1,2,2} \end{pmatrix} \right)$$

Stage II

$$\begin{pmatrix} D_{1,1,1} & D_{1,1,2} \\ D_{1,2,2} & D_{1,2,2} \end{pmatrix} + \begin{pmatrix} D_{2,1,1} & D_{2,1,2} \\ D_{2,2,2} & D_{2,2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

A decomposition induced by a partitioning of D

$$\text{Task 01: } D_{1,1,1} = A_{1,1} B_{1,1}$$

$$\text{Task 02: } D_{2,1,1} = A_{1,2} B_{2,1}$$

$$\text{Task 03: } D_{1,1,2} = A_{1,1} B_{1,2}$$

$$\text{Task 04: } D_{2,1,2} = A_{1,2} B_{2,2}$$

$$\text{Task 05: } D_{1,2,1} = A_{2,1} B_{1,1}$$

$$\text{Task 06: } D_{2,2,1} = A_{2,2} B_{2,1}$$

$$\text{Task 07: } D_{1,2,2} = A_{2,1} B_{1,2}$$

$$\text{Task 08: } D_{2,2,2} = A_{2,2} B_{2,2}$$

$$\text{Task 09: } C_{1,1} = D_{1,1,1} + D_{2,1,1}$$

$$\text{Task 10: } C_{1,2} = D_{1,1,2} + D_{2,1,2}$$

$$\text{Task 11: } C_{2,1} = D_{1,2,1} + D_{2,2,1}$$

$$\text{Task 12: } C_{2,2} = D_{1,2,2} + D_{2,2,2}$$

Figure 3.15 A decomposition of matrix multiplication based on partitioning the intermediate three-dimensional matrix.




Figure 3.16 The task-dependency graph of the decomposition shown in Figure 3.15.

Decomposition Techniques

3. Exploratory Decomposition

- Specially used to decompose the problems having underlying computation *like search-space exploration*.
- Steps:
 1. Partition the *search space into smaller parts*
 2. Search each one of these parts concurrently, *until the desired solutions* are found.

Decomposition Techniques

3. Exploratory Decomposition

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	12

(a)

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	12

(b)

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	12

(c)

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	

(d)

Figure 3.17 A 15-puzzle problem instance showing the initial configuration (a), the final configuration (d), and a sequence of moves leading from the initial to the final configuration.


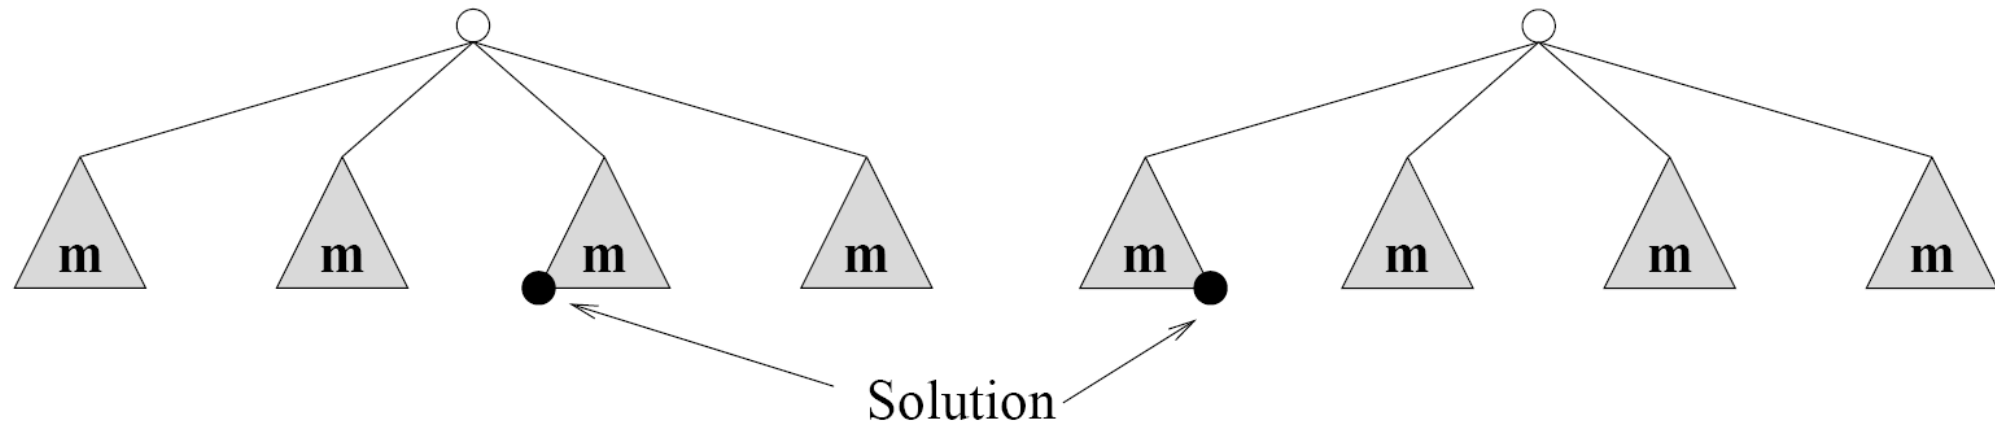


Figure 3.18 The states generated by an instance of the 15-puzzle problem.



Total serial work: $2m+1$

Total parallel work: 1

Total serial work: m

Total parallel work: $4m$

*We may also consider this as 4, (a)
since each task has done 1 step*

(b)

Figure 3.19 An illustration of anomalous speedups resulting from exploratory decomposition.

Decomposition Techniques

4. *Speculative Decomposition*

- Usually used in the problems *where different input values or output of the previous stage* causes many *computationally intensive branches*.
- Speculation is something like *Gamble* or *Risk* or preliminary guess.
- Steps:
 - *Speculate (guess) the output* of previous stage
 - Start performing computations in the *next stage even before the completion of the previous stage*.
 - After the output of the previous stage is available, *if the speculation was correct*, then most of the computation for the next step would have already been done.

Sources

- Slides of Dr. Rana Asif Rahman & Dr. Haroon Mahmood, FAST
- (Chapter 2 & 3) Kumar, V., Grama, A., Gupta, A., & Karypis, G. (1994). Introduction to parallel computing (Vol. 110). Redwood City, CA: Benjamin/Cummings.
- Quinn, M. J. Parallel Programming in C with MPI and OpenMP,(2003).