

Parallel and Distributed Computing

CS3006 (BDS-6A)

Lecture 13

Instructor: Dr. Syed Mohammad Irteza

Assistant Professor, Department of Computer Science, FAST

14 March, 2023

Previous Lecture

- Classic Computing
- Distributed Systems:
 - Motivation
 - Need
 - Examples
 - Benefits and disadvantages
 - Cluster, grid and cloud
- Communication Models (transfer time)
- Message Passing
- C/S and P2P paradigms

Distributed Systems: Issues

- *Programming Complexity*
 - How the various programs communicate together?
 - Which data format on the various network nodes?
 - Need to define (application) protocols
 - Operations synchronization may lead to delay and slowing down
- *Scarce Robustness*
 - Higher chance of errors/faults
- *Hard to Optimize*
 - Lack of a global view

Operating System

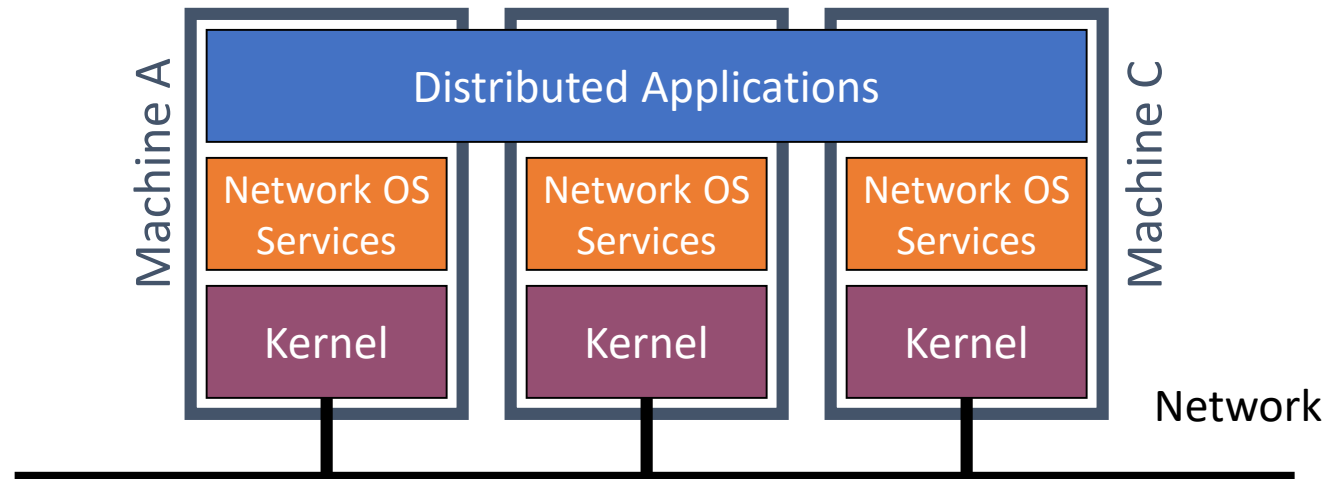
- A program that manages the computer hardware
- Provides a basis for application programs and acts as an intermediary between the user and the hardware
- Offers a reasonable way to solve the problem of creating a usable computing system

Types of Operating Systems:

- Linux and Unix
- Microsoft Windows
- Mac OS X

Network Operating Systems

- Distribution **not hidden** by OS
- But a number of **services** offered to support distribution
 - remote login (ssh), remote copying (scp)
 - distributed file systems (samba, nfs)
- Well adapted to **heterogeneous DS**



Features of Network OS

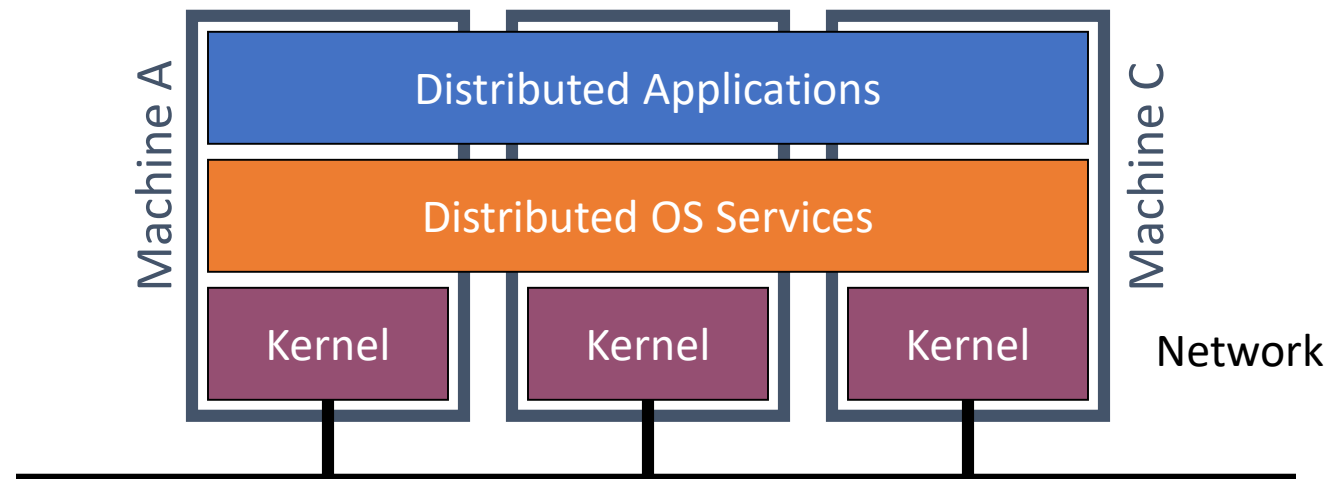
- Provides basic operating system features; *support for processors, protocols, automatic hardware detection*, support *multi-processing of applications*
- Security features; *authentication, authorization, access control*
- Provides files, print, web services, back-up and *replication services*
- Supports Internetworking such as *routing* and *WAN ports*
- User management and support for *login/logoff, remote access, system management*

Network and Operating System Security

- OS: *system must protect itself*
- List of possible breaches is almost endless
- Runaway process could constitute an accidental denial-of-service attack
- Travels over private leased lines, shared lines like the internet, wireless connections, or dial-up lines
- Intercepting these data could be harmful as breaking into a computer
- Interruption of communications could constitute a remote denial-of-service attack

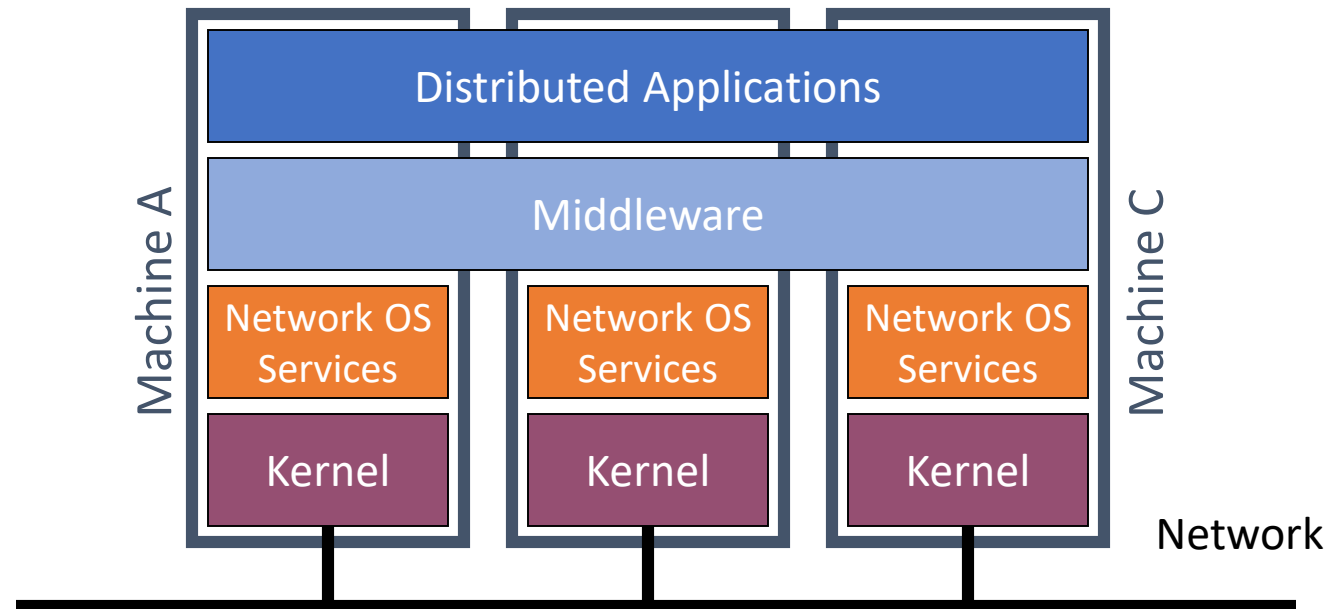
Distributed OSs

- **Completely hide** the distribution: “single system image”
- Each individual node holds a specific software subset of the global aggregate operating system.
- The user has the **illusion** to use one single multiprocessor machine (symmetric multiprocessing, SMP)
- Essentially used for **homogeneous** cluster of machines interconnected with **high performance networks**



Middleware

- Pros and Cons of DOS and NOS
 - DOS are user-friendly but don't scale well. Reverse for NOS
 - How to get the best of both world?
- Solution: **Middleware!**



| System | Description | Main Goal |
|------------------------------------|--|--|
| DOS (Distributed Operating System) | Tightly coupled operating system for multiprocessors and homogenous multicomputers | Hide and manage hardware resources |
| NOS (Network Operating System) | Loosely coupled operating system for heterogeneous multicomputers (LAN and WAN) | Offer local services to remote clients |
| Middleware | Additional layer atop of NOS implementing general-purpose service | Provide distribution transparency |

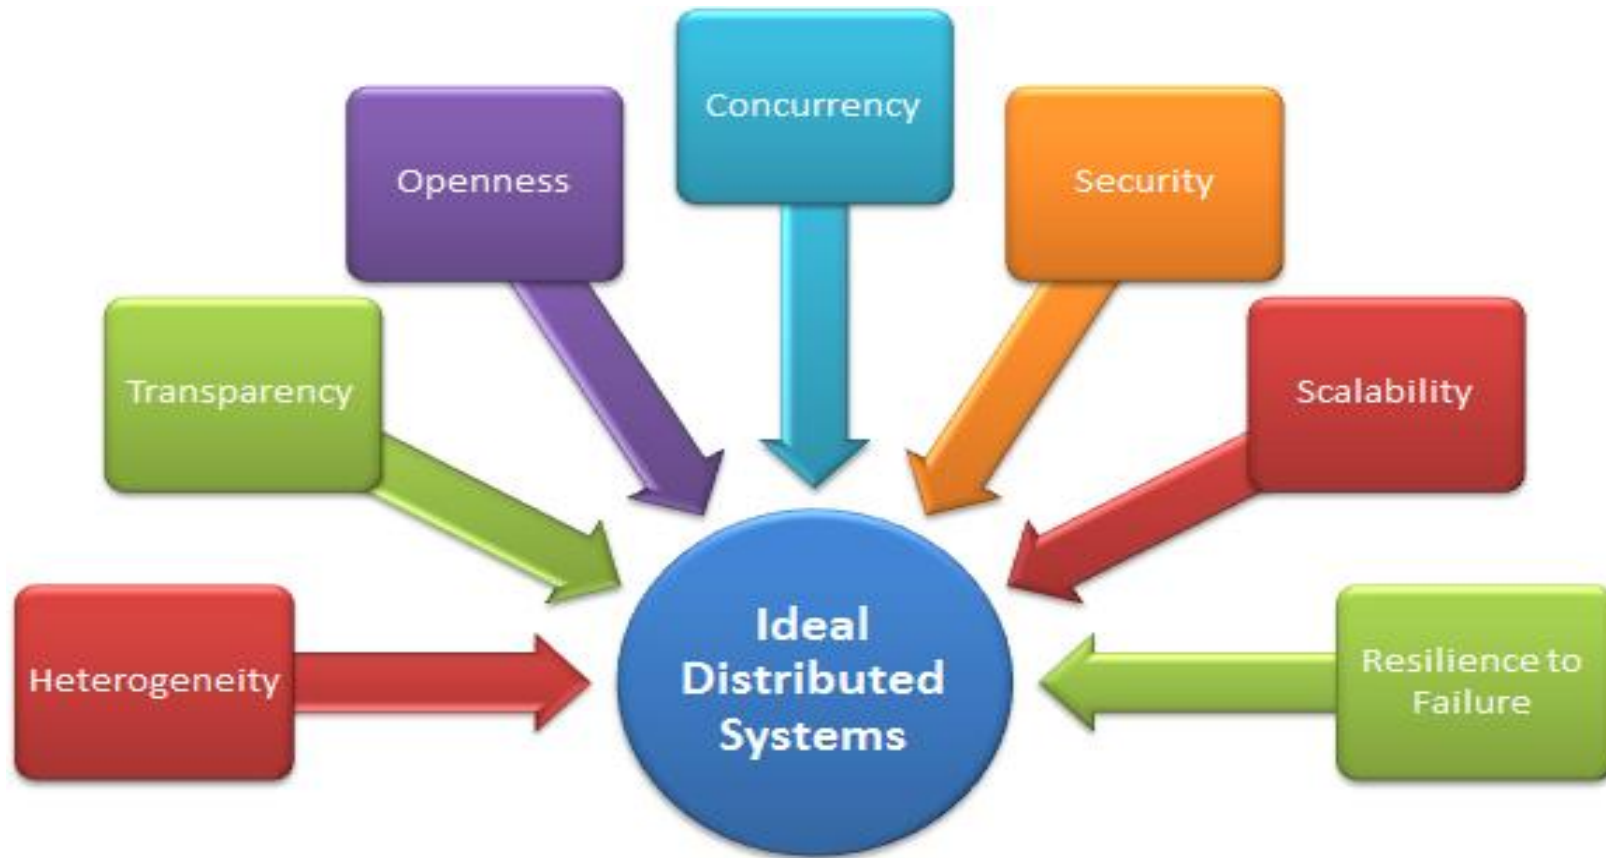
Goals of a Middleware Platform

- Resource sharing
 - The ability to **access and share resources** in a distributed environment
 - The bread and butter of distributed systems
- Transparency
 - The ability to view a distributed system as if it were a single computer
 - Varying **dimensions of transparency** incl. location, access, migration, etc.
 - The **degree of transparency** is a key decision in any systems architecture

Goals of a Middleware Platform

- Openness
 - The offering of services according to standard rules (syntax and semantics)
 - Openness provides support for the key properties of **portability** and **interoperability**
 - Again, the **degree of openness** is a key factor in systems design
- Extensibility
 - The ability to be able to introduce **new or modified functionality**

Challenges of distributed Systems



Heterogeneity

- Distributed systems usually combine *multiple underlying systems* to solve a problem. These systems can differ in the following ways
 - networks
 - computer hardware
 - operating systems
 - programming languages
 - implementations by different developers
- The aim for *middleware* is to provide an abstraction level such that all these heterogeneities *are hidden from the user*

Transparency

- Actual System:
 - a collection of *independent systems*
- User's viewpoint:
 - a *single unified system*
- Transparency is the *concealment of the distributed system from the user*

Challenges for transparency

- To hide the location where a resource is located
- To hide the migration of resource
- To hide the replication of resource
- To hide the differences in data representation
- To hide the failure and recovery of resource
- Hide that a resource may be shared by several competitive users

Openness

- An *open distributed system* is a system that offers services according to standard rules that describe the *syntax* and *semantics* of those services
- Example
 - Implementation of the services in the network
 - Interface Definition
- Different Implementations of the same Interface
 - **Interoperability** - the extent by which *two implementations of systems from different manufacturers* can work together by relying on each other's services as specified by a *common standard*.
 - **Portability** – An application developed for a distributed system *can be executed* (without modification) on a *different system that implements the same interface*

Concurrency

- Any object that represents a *shared resource* in a distributed system must be responsible for ensuring that it *operates correctly* in a concurrent environment
- For an object to be *safe* in a *concurrent environment*, its operations must be *synchronized* in such a way that its data remains *consistent*.

Security

- Resources are shared however they must be kept *secure* – multiple users with their own data
- Components of Security
 - *confidentiality* (protection against disclosure to unauthorized individuals)
 - *integrity* (protection against alteration or corruption)
 - *availability* (protection against interference with the means to access the resources)

Security (Cont..)

- the challenge is to
 - *send sensitive information* in a message over a network in a *secure manner*
 - knowing for sure the *identity of the user* or other agent on whose behalf a message was sent
- *Major Security threats*
 - *Denial of service attacks*
 - *Security of mobile code (Executing Code)*

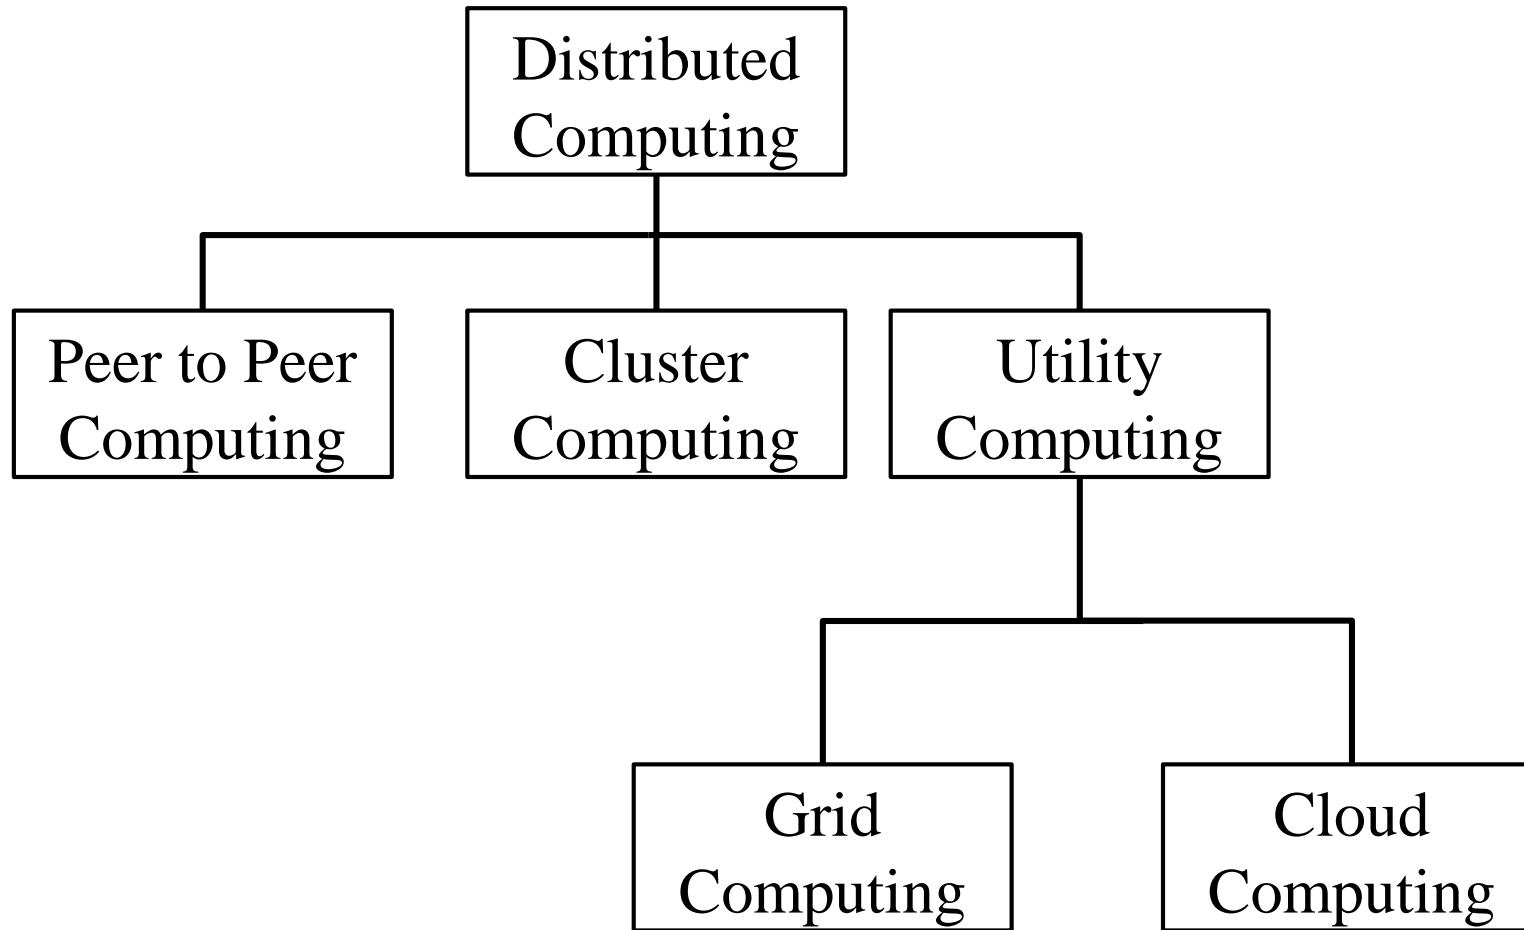
Scalability

- A system is described as *scalable if it will remain* effective when there is a *significant increase in the number of resources and the number of users*
 - *Controlling the cost of physical resources*
 - *Controlling the performance loss*
 - *Preventing software resources running out*
- Dimensions
 - *Scalable with respect to size*
 - *Scalable with respect to geography*
 - *Scalable with respect to administration*
- The challenge is that the system and application software *should not need to change* when the scale of the system increases


Failure Handling

- Failures in a distributed system are partial – that is, some components fail while others continue to function
 - *Masking* Failures - Resending the dropped messages
 - *Tolerating* Failures – Reloading Pages
 - *Recovering* Failures – Roll back – Removing Inconsistency in data
- The design of effective techniques for keeping *replicas* of rapidly changing data up-to-date without excessive loss of performance is a challenge

Distributed Computing




Scalable High Performance Computing




Why we need it?

- Solving grand challenge applications using *modeling, simulation and analysis*




Life Sciences



Aerospace



**Internet &
Ecommerce**



CAD/CAM



Digital Biology



Military Applications

Clusters


- A cluster is a type of a *distributed processing system*, which consists of a collection of interconnected stand-alone computers cooperatively working together as a single, integrated computing resource.
- “stand-alone” (whole computer) that can be used on its own (full hardware and OS).
- Nodes can be used individually or combined

Cluster Computing

- Clusters are usually deployed to improve *speed and/or reliability* over that provided by a *single* computer.
- Much more *cost effective* than single computer of comparable speed or reliability.
- In cluster computing each node within a cluster is an independent system, with its own
 - *operating system*
 - *private memory*
 - in some cases, its *own file system*
- Processors on one machine *cannot directly access the memory* on the other machines
- Programs employ *message passing* to get data and execution code from one machine to another

A cluster is a group of computers connected by a local area network (LAN)

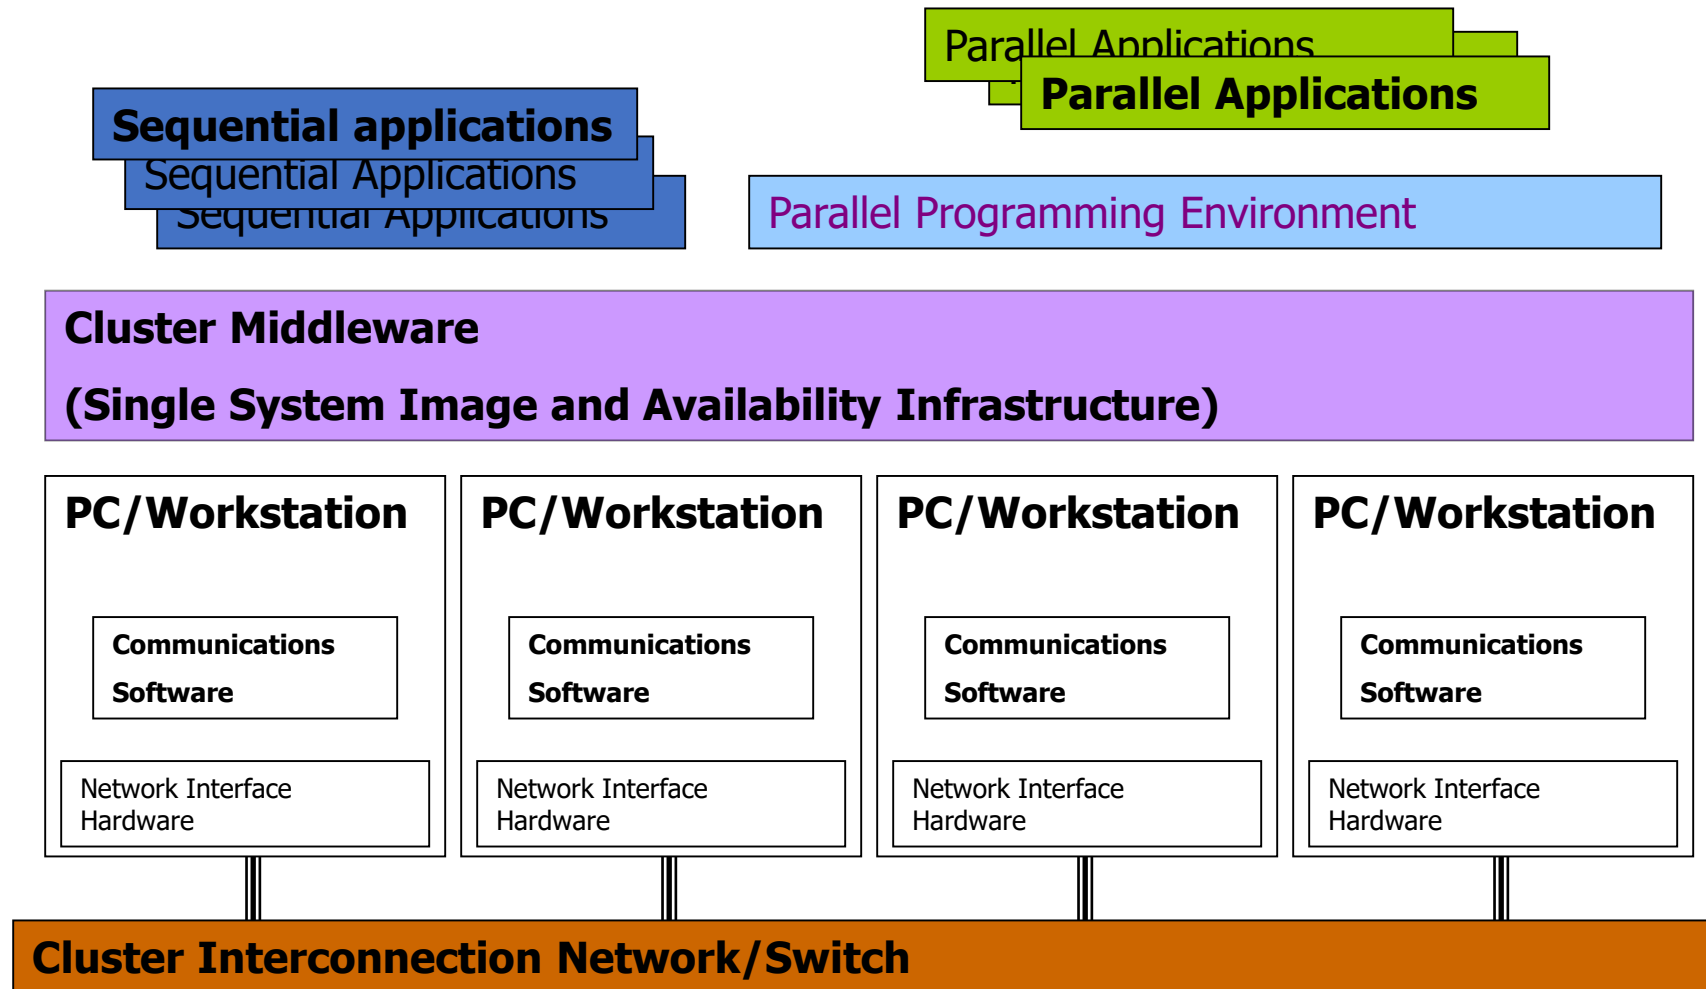
High Performance Cluster



Functioning of a cluster

- A cluster typically has one *head* or *master node* (the main node that handles management tasks) and several *compute nodes*.
- *Management and programming software* is essential to break down *tasks* into *smaller parts* and then *assign* those parts to individual systems.
- The *assigned* systems process their *portions* before they *return* their *results* to cluster management software, which manages and controls the cluster.
- To achieve more with less from a technological standpoint, you can harness the efforts of *numerous smaller and less-expensive processors* to meet or exceed the capabilities of a *larger and more-expensive processor*

Cluster Computer Architecture



Motivations for clustering

- Improving the *time to solution*, especially for *large* and *complex simulations, models, or forecasts*
- Squeezing a better return out of existing investments in computing technology, especially by *increasing overall utilization of computing resources*
- Establishing more *control over computing workloads* and pushing *important jobs to the head of the line* for speedy, reliable processing
- Improving overall computing *reliability, serviceability, and availability*, without necessarily forking over huge amounts of cash for custom-built fault-tolerant systems

Clusters

- A cluster is *tightly coupled*, whereas a grid or a cloud is *loosely coupled*
- All nodes work cooperatively together as a *single integrated computing resource* so conceptually it is smashing up many machines to make a *powerful* machine
- **Types of clusters**
 - High Availability clusters
 - Load balancing clusters
 - High Performance Computing clusters


High Availability Clusters

- These clusters are designed to *provide uninterrupted availability* of data or services (typically web services) to the *end-user community*.
- If a *node fails*, the service can be *restored* without affecting the availability of the services provided by the cluster. While the application will still be available, there will be a *performance drop* due to the missing node.
- The purpose of these clusters is to ensure that a single instance of an application is only ever running on one cluster member at a time but if and when that cluster member is no longer available, the *application will failover to another cluster member*.

High Availability Clusters

- Also known as **Failover Clusters**
- Support server applications that can be *reliably utilized with a minimum of down-time*
- *Redundant computers in groups* that provide continued service even in case of *component failure*
- It is done by *detecting hardware/software faults*, and immediately *restarting the application on another system* without requiring administrative intervention
- High availability clusters implementations are best for *mission-critical applications or databases, mail, file* and *print, web* or *application servers*.

High Availability clusters




Load Balancing Clusters

- Load balancing is required in systems that *handle large volumes of client requests*
- Support *multiuser* and *multitasking* environments
- Overall distribution of the workload of a cluster is *hard to predict* at any particular moment
- Static approach to I/O planning is almost useless
- The two major categories of load balancing are:
 - **Software-based load balancing**
 - consists of special software that is installed on the servers which dispatches requests from clients to servers
 - **Hardware-based load balancing**
 - consists of a specialized switch or router with software to provide load balancing functionality

Load Balancing Cluster

- This type of *cluster distributes incoming requests for resources* or content among multiple nodes running the same programs or having the same content.
- Both the *high availability* and *load-balancing cluster* technologies can be combined to increase the reliability, availability, and scalability of application and data resources that are widely deployed for web, mail, news, or FTP services.
- Every node in the cluster is able to handle requests for the *same content or application*.
- This type of distribution is typically seen in a *web-hosting environment*.

Load Balancing Cluster



High Performance Computing Clusters

- *High performance clusters* are used where
 - *time to solution* is important
 - *problem* is *huge* and can not be executed on one single computer
- The cluster manages the resources needed for the job and assigns the job to a *work queue*
- An ideal solution for problems in which users need to *run many similar jobs* with *different parameters or data sets*
- *Computing is local to a cluster node*, node doesn't communicate with other nodes
- May need *high speed file system access*

High Performance Cluster

- Donald Becker of NASA assembled such a cluster in 1994
- It is also known as the *Beowulf cluster*.
- Essentially, any *group of Linux machines* dedicated to a *single purpose* with a *centralized node* (for coordination) can be called a *Beowulf cluster*.
- These types of clusters *increase performance* and *scalability* for applications, particularly *computationally or data intensive tasks*.
- Applications of such systems include *data mining*, *simulations*, *parallel processing*, *weather modeling*, etc.

References

1. Slides of Dr. Haroon Mahmood
2. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (1994). *Introduction to parallel computing* (Vol. 110). Redwood City, CA: Benjamin/Cummings.
3. Quinn, M. J. *Parallel Programming in C with MPI and OpenMP*, (2003).

Helpful Links:

1. https://en.wikipedia.org/wiki/Java_remote_method_invocation